

Operating Systems and Systems Programming I (*CPS1012*) Short Notes

Aiden Cachia

Contents

1	Intoduction	2
1.1	Operating System Services	2
1.2	System Calls	3
1.3	System Services	5
1.4	Operating System Design	6
2	Processes	9
2.1	Process Scheduling	10

1 Introduction

What is an Operating System?

An Operating System is a program that acts as an intermediary between a user of a computer and the computer hardware.

Operating System Goals

- Execute user programs and make solving user problems easier.
- Make the computer system convenient to use.
- Use the computer hardware in an efficient manner.

Computer System Structure

- Hardware.
- Operating System.
- Application programs.
- Users.

1.1 Operating System Services

- User Interface (CLI, GUI, Touch-screen, batch)
- Program Execution
- I/O Operations
- File-system manipulation
- Communications (On Device, Network)
- Error Detection
- Resource Allocation
- Protection and Security

Command Line Interface

Simple Interface where command is fetched from user, it is executed and the output is displayed.

Graphical User Interface

More user friendly design, making use of;

- Mouse, Keyboard and monitor
- Icons representing files, programs, actions etc ...
- Menus

1.2 System Calls

Ways to interact with the Operating System. They are typically written in a high level language (C or C++) and are used to request a service from the Operating System.

Implementation of System Calls

- A table of all is usually stored in the System-call interface.
- The System-call interface invokes the system-call in OS kernel and returns the status of the system-call and any return values.

System Call Parameters

- Registers
- Block/table
- Stack

Types of System Calls

Process Control

- Create process, terminate process

- End, abort
- Load, execute
- Get process attributes, set process attributes
- Wait for time
- Wait for event, signal event
- Allocate and free memory
- Dump memory
- Debugging
- Locking

File Management

- Create file, delete file
- Open, close
- Read, write, reposition
- Get file attributes, set file attributes

Device Management

- Request device, release device
- Read, write, reposition
- Get device attributes, set device attributes
- Logically attach or detach devices

Information Maintenance

- Get or set
 - time or date
 - system data
 - process, file or device attributes

Communications

- Create, delete communication connection

- Send, receive messages
- Shared memory
- Transfer status information
- Attach or detach remote devices

1.3 System Services

Helps in the development of programs. Some are just system-call interfaces, others more complex.

Types of System Services

- File management
 - Create, delete, copy, rename, print, dump, list and generally manipulate files and directories.
- Status information
 - Basic information (date, time, amount of available memory, disk space, number of users).
 - Detailed logging and Debugging information.
 - Registry to store and retrieve configuration information.
 - File modification
 - Program loading and execution
 - Communications
- Programming language support (Compilers, assemblers, debuggers and interpreters sometimes provided)
- Program loading and execution
- Communications
- Background services
 - Launch at boot
 - Provides facilities like disk checking, process scheduling, error logging, printing.
 - Run in user context not kernel context
 - Known as *services*, *subsystems* or *daemons*
- Application programs

- Unrelated to system
- Run by users
- Not part of OS
- Launched by CLI, Mouse click, finger poke

Linkers and Loaders

A Linker is used to link multiple object files (compiled from source code) into a single executable file.

A Loader is used to load the executable file into memory and start the execution of the program.

Relocation is the process of adjusting the addresses in the object file to match the addresses in memory.

Dynamic linked libraries (DLLs) are libraries that are linked at runtime, rather than at compile time, so they could be used by multiple programs that would require the same version of the library.

Why Applications are Operating System Specific

System calls and file formats are different from one operating system to another, among other things. Because of this applications are not portable between operating systems without modification.

Applications can be made portable:

- By using an interpreted language such as Python or Ruby;
- By using a language which include a VM such as Java;
- By using compiling a program which a User standard language (such as C), then recompiling it in target OSes.

An Application Binary Interface (ABI) is a specification defining requirements for an application on a specific hardware platform.

1.4 Operating System Design

There is no correct way to design an Operating System, but there are some common design principles.

An OS should have User goals (ease of use, convenience, etc ...) and System goals (efficiency, security, etc ...).

Policy and Mechanism

Policy is what will be done, *Mechanism* is how it will be done.

They should be separate, as it would allow for maximum flexibility if policy decisions are to be changed later.

Implementation

Early OSes were written in Assembly, but now they only have the base of the OS written in Assembly, and the rest written in high level languages such as C or C++.

Emulation allows for OS to run on non-native hardware.

The Kernel

The Kernel is the core of the OS, which is used to manage system resources and provides essential services for all other parts of the OS and user-level programs.

Handles tasks such as:

- Process management;
- Memory management;
- Scheduling;
- Interfacing with hardware devices.

Operating System Structure

- Simple Structure - MS-DOS
- More complex - UNIX
- Layered - an abstraction
- Microkernel - Mach

Monolithic - OG Unix

Structure

- System Programs
- The kernel, which contains everything between the system-call interface and the hardware, providing File system, CPU scheduling, memory management, etc ...

Layered

Like an Onion, the Core being Hardware (Layer 0), and the outer being the User Interface (Layer N). Each layer only has access to the layers below it.

Microkernel

A microkernel is the near-minimum amount of software that can provide the mechanisms needed to implement an OS.

- Easier to extend a microkernel
- Easier to port the operating system to new architecture
- More reliable (less code is running in kernel mode)
- More secure

Modules

Most modern OSes make use of Loadable Kernel Modules (LKMs) which can be loaded and unloaded at runtime, Each:

- uses Object-oriented approach;
- core component is separate;
- talks to the other over known interfaces;
- is loadable as needed within the kernel.

Hybrid Systems

Most modern OSes are hybrids of the above systems in order to address performance, security and usability needs.

Dual-mode Operation

Dual-mode operation allows OS to protect itself and other system components.

Mode bit is used to distinguish between user and kernel mode, which is provided by hardware.

2 Processes

A *process* is a program in execution.

Its execution must be sequential. No parallel execution of a single process.

A program becomes a process once it is **active** and **loaded into memory**.

One program can have several processes. Consider multiple users executing the same program.

Process Composition

- *Text Section* - Program code
- *Current Activity* - Program counter, CPU registers
- *Data Section* - Global variables
- *Stack* - Temporary data (function parameters, return addresses, etc ...)
- *Heap* - Dynamically allocated memory

Process State

- *New* - Process is being created
- *Running* - Instructions are being executed
- *Waiting* - Process is waiting for some event to occur
- *Ready* - Process is waiting to be assigned to a processor
- *Terminated* - Process has finished execution

Process Control Block

- *Process State* - New, Ready, Running, Waiting, Terminated
- *Program Counter* - Address of next instruction to be executed
- *CPU Registers* - contents of all process-centric registers
- *CPU Scheduling Information* - Priority, pointers to scheduling queues
- *Memory Management Information* - Page tables, segment tables (memory allocated to the process)
- *Accounting Information* - Amount of CPU and real time used, time limits, account numbers
- *I/O Status Information* - List of I/O devices allocated to the process, list of open files

Threading

TBD

2.1 Process Scheduling

The process scheduler selects which process to run in a way to maximize CPU utilization and throughput.

The Scheduler would have a **ready** queue with all processes in the ready state and a **wait** queue with all processes in the waiting state. The Processes migrate between the two queues.

Context Switch

When the CPU switches to another process, the system must save the state of the old process and load the saved state for the next process via a context switch.

Context-switch time is pure overhead; the system does no useful work while switching. The more complex the OS and the PCB, the longer the context switch time.

The time is dependent on hardware support. Some hardware provides multiple sets of registers per CPU so multiple contexts can be loaded at once.

Process Creation

Process makes process which is his child, and those child create more child, making tree. Process have a unique process ID, and is managed by it.

The options for resource sharing are:

- Child Resources = Parent Resources
- Child Resources \subset Parent Resources
- Child Resources $\not\subset$ Parent Resources

The options for execution are:

- Parent and child execute concurrently
- Parent waits until child terminates

A child process is either a copy of the parent process or a new program.

`fork()` Creates a new process by duplicating the calling process.

`exec()` system call used after a `fork()` to replace the process' memory space with a new program

Parent process calls `wait()` waiting for the child to terminate

Process executes last statement and then asks the operating system to delete it using the `exit()` system call

Parent may terminate the execution of children processing using the `abort()` system call. Some reasons for doing so:

- Child has exceeded allocated resources
- Task assigned to child is not longer required
- The parent is exiting and the OS does not allow a child to continue if its parent terminates

Termination of a tree could be cascading or non-cascading.

- Cascading - Parent process terminates all children
- Non-cascading - Parent process terminates, children continue to be terminated by the OS as it relizes the parent is gone.

If no parent is waiting (did not invoke `wait()`) the process becomes a zombie. If the parent terminated without invoking `wait()`, the process becomes and orphan.